

✓ SMS SPAM DETECTION WITH TENSORFLOW.

IMPORTATION OF THE LIBRARIES

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import tensorflow as tf
6 from tensorflow import keras
7 from tensorflow.keras import layers
```

✓ Loading dataset .read_csv()

```
1 # Using this to load the file from my External Hard drive AND Saving it to Uploaded
2 from google.colab import files
3 uploaded = files.upload()
```

Choose Files spam.csv

- **spam.csv**(text/csv) - 503663 bytes, last modified: 9/20/2019 - 100% done

Saving spam.csv to spam (2).csv

```
1 # when we load the data we use .dropna(axis=1) to drop the unnamed columns with null values.
2 import io
3 # I will use this io.BytesIO(uploaded['spam (2).csv']) to read the csv file from the io that i uploaded the file.
4 df = pd.read_csv(io.BytesIO(uploaded['spam (2).csv']), encoding = 'latin-1').dropna(axis = 1)
5 df.head()
```

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

i will rename the columns v1 and v2 to label and Text respectively


since the target variable is in string form I will encode it numerically using pandas function .map()

```
1 df = df.rename(columns = {'v1': 'label',
2                             'v2' : 'Text'})
3
4 df['label_enc'] = df['label'].map({'ham': 0,
5                                   'spam':1})
6 df.head()
```

	label	Text	label_enc
0	ham	Go until jurong point, crazy.. Available only ...	0
1	ham	Ok lar... Joking wif u oni...	0
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	1
3	ham	U dun say so early hor... U c already then say...	0
4	ham	Nah I don't think he goes to usf, he lives aro...	0

✓ LETS VISUALIZE THE DISTRIBUTION OF HAM AND SPAM DATA

```
1 sns.countplot(x=df['label'])
2 plt.show()
```



✓ NB

we can see that Ham data is comparatively higher than spam data, this is natural, since I will be using embeddings in our deep learning model, I do not have to balance the data. Now let's find the average number of words in all sentences in SMS data

```
1 # Average number of tokens in all sentences
2 avg_words_len = round(sum([len(i.split()) for i in df['Text']]) / len(df['Text']))
3
4 print (avg_words_len)
```

15

✓ Now let's find the Total number of unique words in Corpus

```
1 s = set()
2 for sent in df['Text']:
3     for word in sent.split():
4         s.add(word)
5
6 total_words_length = len(s)
7 print(f"Total Words Length: {total_words_length}")
```

Total Words Length: 15585

✓ SPLIT DATA INTO TRAINING AND TESTING PARTS

```
1 from sklearn.model_selection import train_test_split, cross_val_score
2
3 X, y = np.asarray(df['Text']), np.asarray(df['label_enc'])
4 new_df = pd.DataFrame({'Text': X, 'label': y})
5
6 X_train, X_test, y_train, y_test = train_test_split(new_df['Text'], new_df['label'],
7                                                     test_size = 0.2, random_state = 42)
8
9 X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

((4457,), (4457,), (1115,), (1115,))

✓ Building The Model

first I will build a baseline model and then try to beat the performance of the baseline using deep learning models (*embeddings, LSTM and others*)

I will use **MultinomialNB()**, this performs very well for text Classification when features are discrete like word counts of the words or **tf-idf vectors** (this is a measure that tells us how important or relevant a word is in the document)

```
1 from sklearn.feature_extraction.text import TfidfVectorizer
2 from sklearn.naive_bayes import MultinomialNB
3 from sklearn.metrics import classification_report, accuracy_score, confusion_matrix

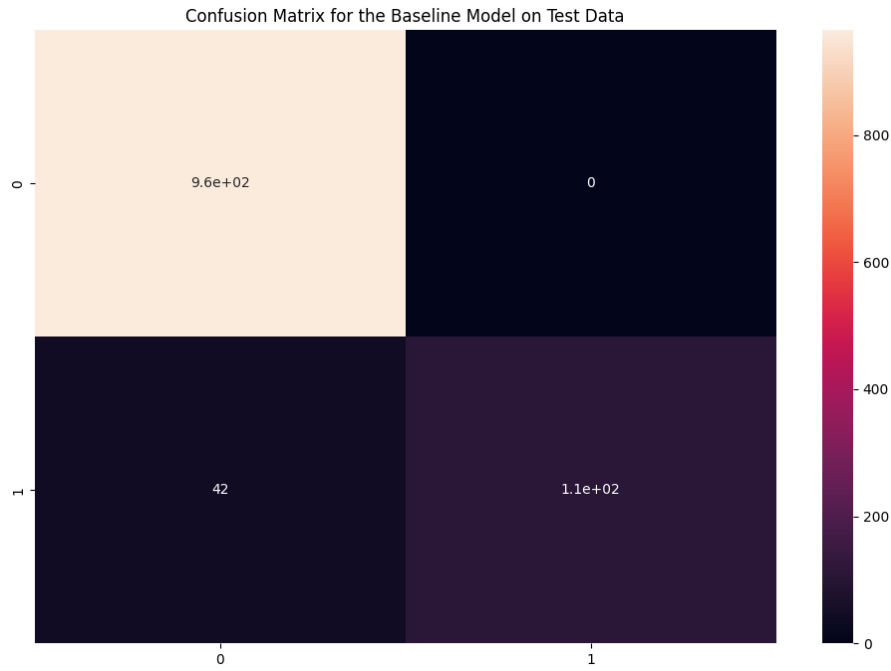
1 tfidf_vec = TfidfVectorizer().fit(X_train)
2
3 X_train_vec, X_test_vec = tfidf_vec.transform(X_train), tfidf_vec.transform(X_test)
4
5
6 baseline_model = MultinomialNB()
7 baseline_model.fit(X_train_vec, y_train)
```

▼ MultinomialNB
MultinomialNB()

▼ Performance of baseline model


```
1 preds = baseline_model.predict(X_test_vec)
2
3 #using the K-fold for cross validation
4
5 def cv_scoring(estimator, X, y):
6     return accuracy_score(y, estimator.predict(X))
7
8
9 print(f"Accuracy on Train Data: {accuracy_score(y_train, baseline_model.predict(X_train_vec))*100}")
10
11 print(f"Accuracy on Test Data : {accuracy_score(y_test, preds)*100}")
12
13 # The Confusion Matrix
14
15 cf_matrix = confusion_matrix(y_test, preds)
16 plt.figure(figsize = (12,8))
17 sns.heatmap(cf_matrix, annot = True)
18 plt.title("Confusion Matrix for the Baseline Model on Test Data")
19 plt.show()
```

Accuracy on Train Data: 97.28516939645502
 Accuracy on Test Data : 96.23318385650225



Model 1: Creating custom Text Vectoriation and embedding layers:

- **Text vectorization** is turning text into numerical representation. example **Bag of words frequency**, **Binary Term frequency**, etc
- **Word Embedding** this is a learned representation of text in which words with related meanings have similar representations. each word assigned a vector



Custom Text Vectorization layer (TensorFlow)

```
1 from tensorflow.keras.layers import TextVectorization
2 MAXTOKENS = total_words_length
3 OUTPUTLEN = avg_words_len
4
5 text_vec = TextVectorization(max_tokens = MAXTOKENS,
6                             standardize = 'lower_and_strip_punctuation',
7                             output_mode = 'int',
8                             output_sequence_length = OUTPUTLEN )
9
10 text_vec.adapt(X_train)
11
12
```

Create an embedding layer

```
1 embedding_layer =layers.Embedding(input_dim = MAXTOKENS,
2                                   output_dim = 128,
3                                   embeddings_initializer = 'uniform',
4                                   input_length = OUTPUTLEN )
5
6
```

input_dim is the size of vocabulary.

output_dim is the dimension of the embedding layer i.e, the size of the vector in which the words will be embedde.

input_length is the length of input sequenc.es

✓ Now lets build and compile model 1 using tensorflow functional API

```
1 input_layer = layers.Input(shape= (1,), dtype = tf.string)
2 vec_layer = text_vec(input_layer)
3 embedding_layer_model = embedding_layer(vec_layer)
4 x = layers.GlobalAveragePooling1D()(embedding_layer_model)
5 x = layers.Flatten()(x)
6 x = layers.Dense(32, activation = 'relu')(x)
7 output_layer = layers.Dense(1, activation = 'sigmoid')(x)
8 model_1 = keras.Model(input_layer, output_layer)
9
10 model_1.compile(optimizer = 'adam', loss= keras.losses.BinaryCrossentropy(label_smoothing = 0.5),
11                                     metrics = ['accuracy'])
12
```

```
1 # The model summary
2 model_1.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 1)]	0
text_vectorization (TextVectorization)	(None, 15)	0
embedding (Embedding)	(None, 15, 128)	1994880
global_average_pooling1d (GlobalAveragePooling1D)	(None, 128)	0
flatten (Flatten)	(None, 128)	0
dense (Dense)	(None, 32)	4128
dense_1 (Dense)	(None, 1)	33
=====		
Total params: 1999041 (7.63 MB)		
Trainable params: 1999041 (7.63 MB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

✓ Callback

lets check and controll the models performance

```
1 from keras.callbacks import EarlyStopping, ReduceLROnPlateau
2
3 es = EarlyStopping(patience = 6,
4                   monitor = 'val_accuracy',
5                   restore_best_weights = True)
6
7 lr = ReduceLROnPlateau(patience = 2,
8                       monitor = 'val_loss',
9                       factor = 0.5,
10                      verbose = 0)
```

✓ Training the model_1

```

1 history = model_1.fit(X_train, y_train, validation_data=(X_test, y_test),
2                       epochs = 20,
3                       batch_size = 32,
4                       callbacks = [lr,es])

```

```

Epoch 1/20
140/140 [=====] - 5s 29ms/step - loss: 0.6060 - accuracy: 0.9053 - val_loss: 0.5775 - val_accuracy: 0.9740 - lr
Epoch 2/20
140/140 [=====] - 5s 34ms/step - loss: 0.5702 - accuracy: 0.9874 - val_loss: 0.5735 - val_accuracy: 0.9803 - lr
Epoch 3/20
140/140 [=====] - 4s 27ms/step - loss: 0.5656 - accuracy: 0.9955 - val_loss: 0.5726 - val_accuracy: 0.9830 - lr
Epoch 4/20
140/140 [=====] - 4s 26ms/step - loss: 0.5641 - accuracy: 0.9982 - val_loss: 0.5724 - val_accuracy: 0.9830 - lr
Epoch 5/20
140/140 [=====] - 5s 32ms/step - loss: 0.5633 - accuracy: 0.9991 - val_loss: 0.5726 - val_accuracy: 0.9803 - lr
Epoch 6/20
140/140 [=====] - 4s 28ms/step - loss: 0.5630 - accuracy: 0.9998 - val_loss: 0.5726 - val_accuracy: 0.9803 - lr
Epoch 7/20
140/140 [=====] - 4s 28ms/step - loss: 0.5628 - accuracy: 0.9998 - val_loss: 0.5725 - val_accuracy: 0.9803 - lr
Epoch 8/20
140/140 [=====] - 5s 33ms/step - loss: 0.5626 - accuracy: 0.9998 - val_loss: 0.5726 - val_accuracy: 0.9803 - lr
Epoch 9/20
140/140 [=====] - 4s 27ms/step - loss: 0.5625 - accuracy: 1.0000 - val_loss: 0.5727 - val_accuracy: 0.9803 - lr

```

Plotting the models history

```


1 print('Train Accuracy: ', np.max(history.history['accuracy'])*100, '%')
2 print('Train Loss: ', np.min(history.history['loss'])*100, '%')
3
4 plt.plot(history.history['accuracy'],label = 'Training Accuracy')
5 plt.plot(history.history['loss'],label = 'Training Loss')
6 plt.title('Model training')
7 plt.ylabel('Accuracy')
8 plt.xlabel('Epoch')
9 plt.legend()
10 plt.show()

```

```

Train Accuracy: 100.0 %
Train Loss: 56.25426173210144 %

```



Helper function for compiling fitting and evaluating the model

```

1 from sklearn.metrics import precision_score, recall_score, f1_score
2
3 def compile_model(model):
4     # Simply compile the model with adam optimizer
5     model.compile(optimizer = keras.optimizers.Adam(),
6                   loss = keras.losses.BinaryCrossentropy(),
7                   metrics = ['accuracy'])
8
9 def fit_model(model, epochs, X_train = X_train, y_train = y_train, X_test = X_test, y_test = y_test):
10     # fit the model with the given epochs, training and test data
11     history = model.fit(X_train,
12                        y_train,
13                        epochs = epochs,
14                        validation_data=(X_test, y_test),
15                        validation_steps=int(0.2*len(X_test)),
16                        callbacks = [lr_scheduler])
17     print('Train Accuracy: ', np.max(history.history['accuracy'])*100, '%')
18     print('Train Loss: ', np.min(history.history['loss'])*100, '%')
19
20     plt.plot(history.history['accuracy'], label = 'Training Accuracy')
21     plt.plot(history.history['loss'], label = 'Training Loss')
22     plt.plot(history.history['val_accuracy'], label = 'validation Accuracy')
23     plt.plot(history.history['val_loss'], label = 'validation Loss')
24     plt.title('Model training')
25     plt.ylabel('Accuracy')
26     plt.xlabel('Epoch')
27     plt.legend()
28     plt.show()
29     return history
30
31 def evaluate_model(model, X,y):
32     # evaluate the model and return accuracy, precision, recall and f1-score
33
34     y_preds = np.round(model.predict(X))
35     accuracy = accuracy_score(y, y_preds)
36     accuracy = accuracy_score(y, y_preds)
37     precision = precision_score(y, y_preds)
38     recall = recall_score(y, y_preds)
39     f1 = f1_score(y, y_preds)
40     print("Model Performance in percentage (%)")
41     model_results_dict = {'accuracy': accuracy * 100,
42                          'precision': precision * 100,
43                          'recall': recall * 100,
44                          'f1-score': f1 * 100}
45
46
47     return model_results_dict

```

✓ Model_2 Bidirectional LSTM

this will effectively improve the networkd accesible information, boosting the context for the algorithm.

```


1 input_layer = layers.Input(shape = (1,), dtype = tf.string)
2 vec_layer = text_vec(input_layer)
3
4 embedding_layer_model = embedding_layer(vec_layer)
5
6 bi_lstm = layers.Bidirectional(layers.LSTM(64,
7                                           activation = 'tanh',
8                                           return_sequences = True))(embedding_layer_model)
9
10 lstm = layers.Bidirectional(layers.LSTM(64))(bi_lstm)
11 flatten = layers.Flatten()(lstm)
12 dropout = layers.Dropout(.1)(flatten)
13
14 x = layers.Dense(32, activation = 'relu')(dropout)
15 output_layer = layers.Dense(1, activation = 'sigmoid')(x)
16
17 model_2 = keras.Model(input_layer, output_layer)
18
19 compile_model(model_2) # Compile the model
20 history_2 = fit_model(model_2, epochs = 20) # fit the model

```

```

Epoch 1/20
140/140 [=====] - 20s 88ms/step - loss: 0.0584 - accuracy: 0.98
Epoch 2/20
140/140 [=====] - 11s 76ms/step - loss: 0.0042 - accuracy: 0.95
Epoch 3/20
140/140 [=====] - 11s 78ms/step - loss: 2.3725e-04 - accuracy:
Epoch 4/20
140/140 [=====] - 13s 90ms/step - loss: 2.9430e-05 - accuracy:
Epoch 5/20
140/140 [=====] - 11s 75ms/step - loss: 1.6673e-05 - accuracy:
Epoch 6/20
140/140 [=====] - 11s 77ms/step - loss: 1.4300e-05 - accuracy:
Epoch 7/20
140/140 [=====] - 11s 79ms/step - loss: 1.2129e-05 - accuracy:
Epoch 8/20
140/140 [=====] - 10s 71ms/step - loss: 9.7068e-06 - accuracy:
Epoch 9/20
140/140 [=====] - 11s 75ms/step - loss: 1.0131e-05 - accuracy:
Train Accuracy: 100.0 %
Train Loss: 0.0009706817763799336 %

```



MODEL_2 EVALUATION

```

1 evaluate_model(model_2,X_test,y_test)

35/35 [=====] - 2s 19ms/step
Model Performance in percentage (%)
{'accuracy': 98.38565022421525,
 'precision': 97.82608695652173,
 'recall': 90.0,
 'f1-score': 93.75}

```

Model- 3 Transfer Learning with USE Encoder

Transfer learning is an approach where one model generated for one job is utilized as the foundation for model on a different task.

USE Layer (universal Sentence Encoder) this converts text into high dimensional vectors that may be used for text categorization, semantic similarity and other language applications

The USE is from Tensorflow_hub and can be used as a layer `.kerasLayer()`

```
1 import tensorflow_hub as hub
```

Model_3 with Sequential api

```
1 model_3 = keras.Sequential()
```


Universal-sentence-encoderlayer

```

1 # directly from tfhub
2 use_layer = hub.KerasLayer("https://tfhub.dev/google/universal-sentence-encoder/4",
3                             trainable = False,
4                             input_shape = [],
5                             dtype = tf.string,
6                             name = 'USE')
7
8 model_3.add(use_layer)
9 model_3.add(layers.Dropout(0.2))
10 model_3.add(layers.Dense(64, activation = keras.activations.relu))
11 model_3.add(layers.Dense(1, activation = keras.activations.sigmoid))
12
13 compile_model(model_3)
14
15 history_3 = fit_model(model_3, epochs =5)


```

```

Epoch 1/5
140/140 [=====] - 10s 48ms/step - loss: 0.2772 - accuracy: 0.92
Epoch 2/5
140/140 [=====] - 4s 29ms/step - loss: 0.0779 - accuracy: 0.977
Epoch 3/5
140/140 [=====] - 6s 46ms/step - loss: 0.0556 - accuracy: 0.983
Epoch 4/5
140/140 [=====] - 4s 30ms/step - loss: 0.0465 - accuracy: 0.986
Epoch 5/5
140/140 [=====] - 4s 30ms/step - loss: 0.0403 - accuracy: 0.988
Train Accuracy: 98.8332986831665 %
Train Loss: 4.025813192129135 %

```

Model training



Now evaluatig all the models

```

1 baseline_model_results = evaluate_model(baseline_model, X_test_vec, y_test)
2 model_1_results = evaluate_model(model_1,X_test,y_test)
3 model_2_results = evaluate_model(model_2,X_test,y_test)
4 model_3_results = evaluate_model(model_3,X_test,y_test)
5
6 total_results = pd.DataFrame({'MultinomialNB Model' : baseline_model_results,
7                               'Custom-Vec-Embedding Model': model_1_results,
8                               'Bidirectional-LSTM Model':model_2_results,
9                               'USE-Transfer Learning Model':model_3_results}).transpose()
10
11 total_results
12

```

```

Model Performance in percentage (%)
35/35 [=====] - 0s 2ms/step
Model Performance in percentage (%)
35/35 [=====] - 0s 11ms/step

```